

Data Engineering Python Coding Challenge

DE120 - J Jatin

Load the CSV File:

- **Explanation:** Before analyzing or manipulating any data, you need to load it into a pandas DataFrame. The CSV file is read using `pandas.read_csv()` which converts the file into a tabular format.

```
PythonCodingChallenge.ipynb X +
[9]: import pandas as pd

# Load the CSV file
df = pd.read_csv('CodingCSV.csv')

# Display the first few rows to check the data
print(df.head())
```

	Year	Industry_aggregation_NZSIOC	Industry_code_NZSIOC	Industry_name_NZSIOC	\
0	2023	Level 1	99999	All industries	
1	2023	Level 1	99999	All industries	
2	2023	Level 1	99999	All industries	
3	2023	Level 1	99999	All industries	
4	2023	Level 1	99999	All industries	

	Units	Variable_code	\
0	Dollars (millions)	H01	
1	Dollars (millions)	H04	
2	Dollars (millions)	H05	
3	Dollars (millions)	H07	
4	Dollars (millions)	H08	

	Variable_name	Variable_category	\
0	Total income	Financial performance	
1	Sales, government funding, grants and subsidies	Financial performance	
2	Interest, dividends and donations	Financial performance	
3	Non-operating income	Financial performance	
4	Total expenditure	Financial performance	

	Value	Industry_code_ANZSIC06
0	930995	ANZSIC06 divisions A-S (excluding classes K633...
1	821630	ANZSIC06 divisions A-S (excluding classes K633...
2	84354	ANZSIC06 divisions A-S (excluding classes K633...
3	25010	ANZSIC06 divisions A-S (excluding classes K633...
4	832964	ANZSIC06 divisions A-S (excluding classes K633...

1. Print Rows of the Data:

- **Explanation:** To get an idea of what your data looks like, you can print the first or last few rows of the DataFrame using `.head()` or `.tail()`.

```
[10]: # Print the first 10 rows
print(df.head(10))
```

	Year	Industry_aggregation_NZSIOC	Industry_code_NZSIOC	Industry_name_NZSIOC	\
0	2023	Level 1	99999	All industries	
1	2023	Level 1	99999	All industries	
2	2023	Level 1	99999	All industries	
3	2023	Level 1	99999	All industries	
4	2023	Level 1	99999	All industries	
5	2023	Level 1	99999	All industries	
6	2023	Level 1	99999	All industries	
7	2023	Level 1	99999	All industries	
8	2023	Level 1	99999	All industries	
9	2023	Level 1	99999	All industries	

	Units	Variable_code	\
0	Dollars (millions)	H01	
1	Dollars (millions)	H04	
2	Dollars (millions)	H05	
3	Dollars (millions)	H07	
4	Dollars (millions)	H08	
5	Dollars (millions)	H09	
6	Dollars (millions)	H10	
7	Dollars (millions)	H11	
8	Dollars (millions)	H12	
9	Dollars (millions)	H13	

	Variable_name	Variable_category	\
0	Total income	Financial performance	
1	Sales, government funding, grants and subsidies	Financial performance	
2	Interest, dividends and donations	Financial performance	
3	Non-operating income	Financial performance	
4	Total expenditure	Financial performance	
5	Interest and donations	Financial performance	
6	Indirect taxes	Financial performance	
7	Depreciation	Financial performance	
8	Salaries and wages paid	Financial performance	
9	Redundancy and severance	Financial performance	

	Value	Industry_code_ANZSIC06
0	930995	ANZSIC06 divisions A-S (excluding classes K633...
1	821630	ANZSIC06 divisions A-S (excluding classes K633...
2	84354	ANZSIC06 divisions A-S (excluding classes K633...
3	25010	ANZSIC06 divisions A-S (excluding classes K633...
4	832964	ANZSIC06 divisions A-S (excluding classes K633...

2. Print the Column Names:

- **Explanation:** DataFrames consist of columns with names. Printing the column names helps you know what data is available and how to reference it.

```
[11]: # Print the column names
      print(df.columns)

Index(['Year', 'Industry_aggregation_NZSIOC', 'Industry_code_NZSIOC',
       'Industry_name_NZSIOC', 'Units', 'Variable_code', 'Variable_name',
       'Variable_category', 'Value', 'Industry_code_ANZSIC06'],
      dtype='object')
```

3. Summary of the DataFrame:

- **Explanation:** The .info() function provides a concise summary, including the data types of columns, non-null counts, and memory usage.

```
[12]: # Summary of the DataFrame
      print(df.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50985 entries, 0 to 50984
Data columns (total 10 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Year                                50985 non-null  int64
 1   Industry_aggregation_NZSIOC         50985 non-null  object
 2   Industry_code_NZSIOC                50985 non-null  object
 3   Industry_name_NZSIOC                50985 non-null  object
 4   Units                              50985 non-null  object
 5   Variable_code                       50985 non-null  object
 6   Variable_name                       50985 non-null  object
 7   Variable_category                   50985 non-null  object
 8   Value                              50985 non-null  object
 9   Industry_code_ANZSIC06              50985 non-null  object
dtypes: int64(1), object(9)
memory usage: 3.9+ MB
None
```

4. Descriptive Statistical Measures:

- **Explanation:** The `.describe()` method provides basic statistical summaries (e.g., mean, standard deviation, minimum, and maximum) for numerical columns.

```
[13]: # Descriptive statistics for numeric columns  
print(df.describe())
```

	Year
count	50985.000000
mean	2018.000000
std	3.162309
min	2013.000000
25%	2015.000000
50%	2018.000000
75%	2021.000000
max	2023.000000

5. Handling Missing Data:

- **Explanation:** Missing data is common in datasets. You can check for missing values and decide whether to fill them (e.g., with 0 or a mean value) or drop rows with missing data.

```
[14]: # Check for missing data  
print(df.isnull().sum())  
  
# Example: Fill missing values with a specific value (e.g., 0)  
df.fillna(0, inplace=True)  
  
# Alternatively, you can drop rows with missing values  
df.dropna(inplace=True)
```

Year	0
Industry_aggregation_NZSIOC	0
Industry_code_NZSIOC	0
Industry_name_NZSIOC	0
Units	0
Variable_code	0
Variable_name	0
Variable_category	0
Value	0
Industry_code_ANZSIC06	0

dtype: int64

(No Missing Data in this Dataset)

6. Sorting the DataFrame:

- **Explanation:** Sorting helps you reorganize the dataset based on one or more columns, such as sorting by 'Value' in ascending or descending order.

```
[15]: # Sort by the 'Value' column in descending order
df_sorted = df.sort_values(by='Value', ascending=False)
print(df_sorted.head())
```

	Year	Industry_aggregation_NZSIOC	Industry_code_NZSIOC	\
40287	2015	Level 4	KK112	
40219	2015	Level 3	KK11	
22078	2019	Level 4	LL122	
31348	2017	Level 4	LL122	
31347	2017	Level 4	LL122	

	Industry_name_NZSIOC	Units	Variable_code	\
40287	Financial Asset Investing	Dollars (millions)	H26	
40219	Finance	Dollars (millions)	H26	
22078	Non-Residential Property Operation	Dollars (millions)	H28	
31348	Non-Residential Property Operation	Dollars (millions)	H28	
31347	Non-Residential Property Operation	Dollars (millions)	H27	

	Variable_name	Variable_category	Value	\
40287	Fixed tangible assets	Financial position	S	
40219	Fixed tangible assets	Financial position	S	
22078	Disposals of fixed assets	Financial position	S	
31348	Disposals of fixed assets	Financial position	S	
31347	Additions to fixed assets	Financial position	S	

	Industry_code_ANZSIC06
40287	ANZSIC06 group K624
40219	ANZSIC06 groups K621, K622, K623, and K624
22078	ANZSIC06 class L671200
31348	ANZSIC06 class L671200
31347	ANZSIC06 class L671200

9. Using Lambda Functions:

- **Explanation:** Lambda functions allow you to apply quick, anonymous functions without explicitly defining them. Useful for simple operations.

pd.to_numeric(df['Value'], errors='coerce'): This function converts the 'Value' column to numeric data. Any invalid string values (like letters or symbols) will be converted to NaN (Not a Number).

Lambda function: The lambda function will now correctly compare numeric values in the 'Value' column to 500,000 and categorize them as "High" or "Low."

```
[15]: # Convert the 'Value' column to numeric, setting errors='coerce' to handle non-numeric values
df['Value'] = pd.to_numeric(df['Value'], errors='coerce')

# Now apply the lambda function to create a new column 'Value_Category'
df['Value_Category'] = df['Value'].apply(lambda x: 'High' if x > 500000 else 'Low')

# Print first few rows to check the new column
print(df[['Value', 'Value_Category']].head())
```

	Value	Value_Category
0	930995.0	High
1	821630.0	High
2	84354.0	Low
3	25010.0	Low
4	832964.0	High

11. Number of Columns in the Dataset:

- **Explanation:** The shape of the DataFrame gives you the number of rows and columns. `shape[1]` returns the number of columns.

```
[9]: # Get the number of columns
print(f"Number of columns: {df.shape[1]}")

Number of columns: 10
```

13. How the Dataset is Indexed:

- **Explanation:** DataFrames are indexed (usually by row numbers), but the index can also be custom (e.g., dates). Checking the index tells you how data is organized.

```
[10]: # Check dataset indexing
print(df.index)

Int64Index([ 0, 1, 2, 3, 4, 5, 6, 7, 8,
            9,
            ...,
            50975, 50976, 50977, 50978, 50979, 50980, 50981, 50982, 50983,
            50984],
            dtype='int64', length=50985)
```

14. Number of Observations in the Dataset:

- **Explanation:** The number of rows or "observations" gives you an idea of the dataset's size. This can be checked using `shape[0]`.

```
[11]: # Get the number of rows
print(f"Number of observations: {df.shape[0]}")

Number of observations: 50985
```

13. Visualizing DataFrame:

Explanation: The code processes a CSV file using Pandas to prepare data for a scatter plot. It reads the file, converts the Value column to numeric (replacing invalid entries with NaN), and drops rows with missing values. It samples 1000 rows randomly. Using Matplotlib, it creates a scatter plot of Year (x-axis) versus Value (y-axis), with transparency and grid lines to improve readability. This approach ensures efficient and clear visualization of large datasets.

```
[17]: import pandas as pd
import matplotlib.pyplot as plt

# Load the CSV file
file_path = 'CodingCSV.csv'
df = pd.read_csv(file_path)

# Convert the 'Value' column to numeric, replacing non-numeric entries with NaN
df['Value'] = pd.to_numeric(df['Value'], errors='coerce')

# Drop rows with NaN in 'Value' or other relevant columns
df_clean = df.dropna(subset=['Value'])

# Sample a subset of the data to reduce size for plotting
sampled_df = df_clean.sample(n=1000, random_state=42)

# Scatter plot: Year vs. Value
plt.figure(figsize=(8, 6))
plt.scatter(sampled_df['Year'], sampled_df['Value'], alpha=0.6, s=10)
plt.title('Scatter Plot: Year vs Value', fontsize=14)
plt.xlabel('Year', fontsize=12)
plt.ylabel('Value (Dollars)', fontsize=12)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

